


MINE-1 Knowledge Retention — KBExtractor vs KGGen

same extraction backbone, same judge



Judge: deepseek-r1:32b (32B params, on-prem vLLM, temp 0, CoT) | Extraction core (both systems): deepseek-r1:32b (DeepSeek-R1 distilled Qwen-32B, 32B params, temp 0)
Retriever: all-MiniLM-L6-v2 (22.7M params), top-k=8 nearest nodes + 2-hop expansion | MINE-1 = % of 15 known facts entailed by the retrieved sub-graph, per article