


# Judge-LLM ablation — KBExtractor > KGGen under every judge



Each judge scores identical graphs (same retriever + metric); bars are over that judge's own article set (n). GPT-5 capped at 10 articles for cost ( $\approx$  \$0.17/article). Absolute level varies with judge strictness; the gap persists.